

CMPS 224/396AA: GPU COMPUTING
ASSIGNMENT 1

In this assignment, you will write a CUDA program that computes the elementwise maximum of two vectors of double precision values. That is, you will compute a vector `c` from two vectors `a` and `b` such that `c[i]` is the maximum of `a[i]` and `b[i]` for all `i`.

Instructions

1. Place the files provided with this assignment in a single directory. The files are:
 - `main.cu`: contains setup and sequential code
 - `kernel.cu`: where you will implement your code (you should only modify this file)
 - `common.h`: for shared declarations across `main.cu` and `kernel.cu`
 - `timer.h`: to assist with timing
 - `Makefile`: used for compilation
2. Edit `kernel.cu` where `TODO` is indicated to implement the following:
 - Allocate device memory
 - Copy data from the host to the device
 - Configure and invoke the CUDA kernel
 - Copy the results from the device to the host
 - Free device memory
 - Perform the computation in the kernel
3. Compile your code by running: `make`
4. Test your code by running: `./vecmax`
 - If you are using the university cluster, do not forget to use the submission system. Do not run on the head node!
 - For testing on different input sizes, you can provide your own value for the number of vector elements: `./vecmax <M>` (example: `./vecmax 1000000`)

Submission

Submit your modified `kernel.cu` file via Moodle by the due date. Do not submit any other files or compressed folders.